

Assignment-3

AIM: Locate dataset (e.g., sample_weather.txt) for working on weather data which reads the text input files and finds average for temperature, dew point and wind speed.

Let us summarize how Hadoop works step by step:

- Input data is broken into blocks of size 128 Mb and then blocks are moved to different nodes.
- Once all the blocks of the data are stored on data-nodes, the user can process the data.
- Resource Manager then schedules the program (submitted by the user) on individual nodes.
- Once all the nodes process the data, the output is written back to HDFS.

Weather Analyse (average Temperature, dew point, wind speed)

Weather.java

```
import java.io.IOException;
import java.util.ArrayList;
import java.util.Iterator;
import java.util.List;
import java.util.StringTokenizer;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.conf.Configured;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.FileInputFormat;
import org.apache.hadoop.mapred.FileOutputFormat;
import org.apache.hadoop.mapred.JobClient;
import org.apache.hadoop.mapred.JobConf;
import org.apache.hadoop.mapred.KeyValueTextInputFormat;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reducer;
import org.apache.hadoop.mapred.Reporter;
import org.apache.hadoop.util.*;
/**
 * This is an Hadoop Map/Reduce application for Working on weather data It reads
 * the text input files, breaks each line into stations weather data and finds
 * average for temperature, dew point , wind speed. The output is a locally
 * sorted list of stations and its 12 attribute vector of average temp , dew ,
 * wind speed of 4 sections for each month.
 */
public class Weather extends Configured implements Tool {
    final long DEFAULT_SPLIT_SIZE = 128 * 1024 * 1024;
    /**
     * Map Class for Job 1
     * For each line of input, emits key value pair with
     * station_yearmonth_sectionno as key and 3 attribute vector with
     * temperature , dew point , wind speed as value. Map method will strip the
     * day and hour from field and replace it with section no (
     * <b>station_yearmonth_sectionno</b>, <b>temperature,dew point , wind
```

```

* speed></b>).
*/
public static class MapClass extends MapReduceBase
    implements Mapper<LongWritable, Text, Text, Text> {
private Text word = new Text();
private Text values = new Text();
public void map(LongWritable key, Text value,
                OutputCollector<Text, Text> output,
                Reporter reporter) throws IOException {
    String line = value.toString();
    StringTokenizer itr = new StringTokenizer(line);
    int counter = 0;
    String key_out = null;
    String value_str = null;
    boolean skip = false;
    loop:while (itr.hasMoreTokens() && counter<13) {
        String str = itr.nextToken();
        switch (counter) {
            case 0:
                key_out = str;
                if(str.contains("STN")){//Ignoring rows where stationid is all 9
                    skip = true;
                    break loop;
                }else{
                    break;
                }
            case 2:
                int hour = Integer.valueOf(str.substring(str.lastIndexOf("_")+1, str.length()));
                str = str.substring(4,str.lastIndexOf("_")-2);
                if(hour>4 && hour<=10){
                    str = str.concat("_section1");
                }else if(hour>10 && hour<=16){
                    str = str.concat("_section2");
                }else if(hour>16 && hour<=22){
                    str = str.concat("_section3");
                } else{ str = str.concat("_section4");
                }
                key_out = key_out.concat("_").concat(str);
                break;
            case 3://Temperature
                if(str.equals("9999.9")){//Ignoring rows temperature is all 9
                    skip = true;
                    break loop;
                }else{
                    value_str = str.concat(" ");
                    break;
                }
            case 4://Dew point
                if(str.equals("9999.9")){//Ignoring rows where dewpoint all 9
                    skip = true;
                    break loop;
                }
        }
    }
}

```

```

        }else{
            value_str = value_str.concat(str).concat(" ");
            break;
        }
    case 12://Wind speed
        if(str.equals("999.9")){//Ignoring rows wind speed is all 9
            skip = true;
            break loop;
        }else{ value_str = value_str.concat(str).concat(" ");
            break;
        }
    default: break;
}
counter++;
}

if(!skip){
    word.set(key_out);
    values.set(value_str);
    output.collect(word, values);
}
}

}

/**
 * Reducer Class for Job 1
 * A reducer class that just emits 3 attribute vector with average
 * temperature , dew point , wind speed for each of the section of the month for each input
 */
public static class Reduce extends MapReduceBase
    implements Reducer<Text, Text, Text, Text>{
private Text value_out_text = new Text();
public void reduce(Text key, Iterator<Text> values,
                  OutputCollector<Text, Text> output, Reporter reporter) throws IOException {
    double sum_temp = 0;
    double sum_dew = 0;
    double sum_wind = 0;
    int count = 0;

    while (values.hasNext()) {
        String str = values.next().toString();
        StringTokenizer itr = new StringTokenizer(str);
        int count_vector = 0;
        while (itr.hasMoreTokens()) {
            String nextToken = itr.nextToken(" ");
            if(count_vector==0){
                sum_temp += Double.valueOf(nextToken);
            }
            if(count_vector==1){
                sum_dew += Double.valueOf(nextToken);
            }
            if(count_vector==2){
                sum_wind += Double.valueOf(nextToken);
            }
        }
        output.collect(key, new Text(sum_temp + " " + sum_dew + " " + sum_wind));
    }
}
}

```

```

        count_vector++;
    }
    count++;
}
double avg_tmp = sum_temp / count;
double avg_dew = sum_dew / count;
double avg_wind = sum_wind / count;
System.out.println(key.toString()+" count is "+count+" sum of temp is "+sum_temp+" sum of
dew is "+sum_dew+" sum of wind is "+sum_wind+"\n");
String value_out = String.valueOf(avg_tmp).concat(
").concat(String.valueOf(avg_dew)).concat(" ").concat(String.valueOf(avg_wind));
value_out_text.set(value_out);
output.collect(key, value_out_text);
}
}

static int printUsage() {
    System.out.println("weather [-m <maps>] [-r <reduces>] <job_1 input> <job_1 output> <job_2
output>");
    ToolRunner.printGenericCommandUsage(System.out);
    return -1;
}
/**
 * The main driver for weather map/reduce program.
 * Invoke this method to submit the map/reduce job.
 * @throws IOException When there is communication problems with the job tracker.
 */
public int run(String[] args) throws Exception {
    Configuration config = getConf();
    // We need to lower input block size by factor of two
    JobConf conf = new JobConf(config, Weather.class);
    conf.setJobName("Weather Job1");

    // the keys are words (strings)
    conf.setOutputKeyClass(Text.class);
    // the values are counts (ints)
    conf.setOutputValueClass(Text.class);

    conf.setMapOutputKeyClass(Text.class);
    conf.setMapOutputValueClass(Text.class);
    conf.setMapperClass(MapClass.class);
    //conf.setCombinerClass(Combiner.class);
    conf.setReducerClass(Reduce.class);
    List<String> other_args = new ArrayList<String>();
    for(int i=0; i < args.length; ++i) {
        try {
            if ("-m".equals(args[i])) {
                conf.setNumMapTasks(Integer.parseInt(args[++i]));
            } else if ("-r".equals(args[i])) {
                conf.setNumReduceTasks(Integer.parseInt(args[++i]));
            } else {
                other_args.add(args[i]);
            }
        }
    }
}
```

```

        } catch (NumberFormatException except) {
            System.out.println("ERROR: Integer expected instead of " + args[i]);
            return printUsage();
        } catch (ArrayIndexOutOfBoundsException except) {
            System.out.println("ERROR: Required parameter missing from " +
                args[i-1]);
            return printUsage();
        }
    }

    // Make sure there are exactly 2 parameters left.
    FileInputFormat.setInputPaths(conf, other_args.get(0));
    FileOutputFormat.setOutputPath(conf, new Path(other_args.get(1)));
    JobClient.runJob(conf);
    return 0;
}

public static void main(String[] args) throws Exception {
    int res = ToolRunner.run(new Configuration(), new Weather(), args);
    System.exit(res);
}
}
}

```

Input: sample_weather.txt (sample)

690190 13910 20060201_0	51.75	33.024	1006.324	943.924	15.024	10.724	22.0	28.9	0.001999.9	000000
690190 13910 20060201_1	54.74	33.024	1006.324	943.924	15.024	10.724	22.0	28.9	0.001999.9	00000
690190 13910 20060201_2	50.59	33.024	1006.324	943.924	15.024	10.724	22.0	28.9	0.001999.9	00000
690190 13910 20060201_3	51.67	33.024	1006.324	943.924	15.024	10.724	22.0	28.9	0.001999.9	000000
690190 13910 20060201_4	65.67	33.024	1006.324	943.924	15.024	10.724	22.0	28.9	0.001999.9	000000
690190 13910 20060201_5	55.37	33.024	1006.324	943.924	15.024	10.724	22.0	28.9	0.001999.9	000000
690190 13910 20060201_6	49.26	33.024	1006.324	943.924	15.024	10.724	22.0	28.9	0.001999.9	000000
690190 13910 20060201_7	55.44	33.024	1006.324	943.924	15.024	10.724	22.0	28.9	0.001999.9	000000
690190 13910 20060201_8	64.05	33.024	1006.324	943.924	15.024	10.724	22.0	28.9	0.001999.9	000000

Output: part-00000.txt (on Hadoop)

690190_02_section1	53.87166666666666	25.899999999999995	7.774999999999998
690190_02_section2	54.76125000000001	25.900000000000006	7.774999999999999
690190_02_section3	53.2504166666667	25.899999999999995	7.774999999999996
690190_02_section4	52.4470833333333	25.900000000000006	7.774999999999999

Weather Data Analysis Steps to run:

26. Starting Hadoop

\$ start-all.sh

27. Made A folder “WeatherAssi” and write Weather.java code.

28. Create new folder for input data.

29. Add input text file in the input data folder.

30. Create new folder to hold java class files.

31. Set HADOOP_CLASSPATH environment variable.

\$ export HADOOP_CLASSPATH=\$(hadoop classpath)

32. Create a directory on HDFS

```
$ hdfs dfs -mkdir /WeatherTut  
$ hdfs dfs -mkdir /WeatherTut/Input
```

33. Checking on localhost:9870

34. Upload the input file (device) to that directory.

```
$ hdfs dfs -put input_data/sample_wheater.txt /WeatherTut/Input
```

35. Compile the java code:

```
$ javac -classpath $(HADOOP_CLASSPATH) -d '/home/huser/Desktop/WeatherAssi/  
weather_java' /home/huser/Desktop/WeatherAssi/Weather.java
```

36. Creation .jar file of classes:

```
$ jar -cvf weathertut.jar -C /home/huser/Desktop/WeatherAssi/weather_java/ .
```



37. Running the jar file on Hadoop

```
$ hadoop jar weather.jar Weather /WeatherTut/Input /WeatherTut/Output
```



38. Check output on localhost:9870 /localhost:50070

Browse Directory

/WeatherTut/Output

Show 25 entries Search:

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rw-r--r--	huser	supergroup	0 B	Apr 14 23:48	1	128 MB	_SUCCESS
-rw-r--r--	huser	supergroup	296 B	Apr 14 23:48	1	128 MB	part-00000

Showing 1 to 2 of 2 entries Previous 1 Next

File information - part-00000

Download Head the file (first 32K) Tail the file (last 32K)

Block information -- Block 0

Block ID: 1073741918
Block Pool ID: BP-1388353168-127.0.1.1-1647528100285
Generation Stamp: 1094
Size: 296
Availability:
• ubuntu-college

File contents

690190_02_section1	53.87166666666666 25.89999999999995 7.774999999999998
690190_02_section2	54.76125000000001 25.900000000000006 7.774999999999999
690190_02_section3	53.25041666666667 25.89999999999995 7.774999999999996
690190_02_section4	52.44708333333333 25.900000000000006 7.774999999999999

Close

39. Stop Hadoop services:

\$ stop-all.sh